

Advanced Management of Data

Exercise 3 Topic 3:

Object-Relational Database Systems

Arrays

- Task: create new types `multiemailtype` and `multitelephonenumber`, that can store more than one entry and find a way to enforce the constraints of `emailtype` on the entries of `multiemailtype`

Array Types

```
CREATE FUNCTION enforceOneAtInTheMiddle(emails VARCHAR[]) RETURNS BOOLEAN AS $$  
DECLARE  
    email VARCHAR;  
BEGIN  
    FOREACH email IN ARRAY emails LOOP  
        IF email !~ '^[^@]+@[^@]+$' THEN  
            RETURN FALSE;  
        END IF;  
    END LOOP;  
    RETURN TRUE;  
END;  
$$ LANGUAGE plpgsql;  
  
CREATE DOMAIN multiemailtype AS VARCHAR[] CHECK(enforceOneAtInTheMiddle(value));  
  
CREATE DOMAIN multitelephonetype AS VARCHAR[];
```

- Task: exchange the old types with the new multi-types in persons and add +49 234 56789 and Max@example.com to Max Mustermann

Arrays Appended

```
ALTER TABLE persons
  ALTER email TYPE multiemailtype
    USING ARRAY[email]::multiemailtype,
  ALTER telephone TYPE multitelephonetype
    USING ARRAY[telephone]::multitelephonetype;
```

- now you could add new email and telephone like

```
UPDATE persons
  SET email = email || 'Max@example.com'::VARCHAR,          -- use simple array concatenation
      telephone = array_append(telephone, '+49 234 56789') -- or a function doing the same
  WHERE name = ('Max', 'Mustermann')::nametype;
```

- as a reminder: as the table professors is inherited from persons, the columns are changed there, too

Arrays

NF1 ↔ NF2

- consider the following table in First Normal Form that stores information about books

```
CREATE TABLE book_1nf (title VARCHAR, author VARCHAR, year INT, month VARCHAR, day INT, keyword VARCHAR);
```

```
INSERT INTO book_1nf VALUES
```

```
  ('Selling', 'Stein', 2009, 'April', 1, 'Profit'), ('Selling', 'Stein', 2009, 'April', 1, 'Strategy'),
  ('Selling', 'Jahn', 2009, 'April', 1, 'Profit'), ('Selling', 'Jahn', 2009, 'April', 1, 'Strategy'),
  ('Report', 'Jahn', 2017, 'June', 14, 'Profit'), ('Report', 'Jahn', 2017, 'June', 14, 'Staff'),
  ('Report', 'Frey', 2017, 'June', 14, 'Profit'), ('Report', 'Frey', 2017, 'June', 14, 'Staff');
```

- as you can see, there is much redundancy as the title and the date information is copied for each combination of author and keyword
- therefore we are going to drop the First Normal Form and storing the information in the Non First Normal Form
- Task: create a new type `date_type`, that aggregates the presented date information into just one field, and also create a new table `book_nf2`, that makes use of this new type `date_type` and also uses lists of authors and keywords to prevent the redundant information in the presented table `book_1nf`
- finally, copy all information from `book_1nf` to `book_nf2` by nesting the information into just two datasets and avoid duplicates

Arrays

Nest from 1NF to NF²

```
CREATE TYPE datatype AS (year INTEGER, month VARCHAR, day INTEGER);
```

```
CREATE TABLE book_nf2  
  (title VARCHAR, authors VARCHAR[], dateymd datatype, keywords VARCHAR[]);
```

```
INSERT INTO book_nf2  
  SELECT title,  
         array_agg(DISTINCT author),           -- DISTINCT avoids duplicate entries  
         (year, month, day)::datatype as dateymd, -- the DBMS cannot do this cast on its own  
         array_agg(DISTINCT keyword)  
  FROM book_1nf GROUP BY title, dateymd;
```

- now, you know the way from 1NF to NF², but you also should know the other way around
- Task: copy the information back from book_nf2 to book_1nf by unnesting the data

Arrays

Unnest from NF² to 1NF

```
INSERT INTO book_1nf
  SELECT title,
         unnest(authors),
         (dateymd).year,
         (dateymd).month,
         (dateymd).day,
         keyword
FROM (
  SELECT title,
         authors,
         dateymd,
         unnest(keywords) as keyword -- unnest keywords first in sub-query
  FROM book_nf2
) as unnestfirstlevel; -- the sub-query has to be named but the name doesn't matter
```

-- unnest authors and dateymd second in main-query

Object Identifiers (OIDs)

- OIDs are used internally by PostgreSQL as primary keys for various system tables
- they are not added to user-created tables by default but this can be explicitly demanded

```
CREATE TABLE tablename (tabledefinition) WITH OIDS;
```

- as OIDs are currently implemented as unsigned 4-byte integer, this will not suffice for database-wide uniqueness in large databases or even large tables and their use as primary keys is discouraged
- they are best used only for references to system tables
- real references/pointers are not supported